

Conference Contribution: Geofinitism — A Geometric-Finitist Philosophy for the Age of AI

Author: Kevin R. Haylett

Presented at: [Conference Name], 2026

Format: 20-minute presentation + Q&A with LLM participants

SUMMARY FOR CONFERENCE PROCEEDINGS

Geofinitism: Meaning as Finite Geometry

Kevin R. Haylett

Abstract: Geofinitism proposes a foundational shift in how we understand meaning, knowledge, and reality. Starting from the primacy of finite measurement, it argues that all we can know exists as trajectories in a high-dimensional geometric manifold: the Grand Corpus of all symbolic transductions. The framework rests on Five Pillars: (1) Geometric Container Space—meaning is position and motion in manifolds; (2) Approximations and Measurements—all symbols are lossy compressions of unmeasurable reality; (3) Dynamic Flow—meaning flows nonlinearly through cascading layers; (4) Useful Fiction—frameworks validate through utility, not correspondence; (5) Finite Reality—all measurements are bounded; infinities are demoted to symbolic procedures. Geofinitism finds empirical support in LLM behavior: degradation experiments reveal structured attractors (paranoia, recursion, Zen-like paradoxes) rather than random noise; transformer attention mechanisms are reinterpreted as implicit Takens embeddings that reconstruct semantic manifolds from token sequences. The Takens-Based Transformer (MARINA) operationalizes this view, replacing quadratic attention with explicit delay-coordinate reconstruction, achieving linear complexity and fixed memory while revealing task-specific geometric structures. Geofinitism offers practical tools for AI alignment (basin engineering), interpretability (attractor mapping), and security (embedding integrity). It resonates with relational ontologies across traditions while providing a scientific, measurement-grounded alternative to infinite-accommodating frameworks.

Keywords: Geofinitism, finite measurement, nonlinear dynamics, attractor basins, Takens' theorem, LLM, Grand Corpus, useful fiction

PRESENTATION SLIDES

SLIDE 1: Title

Geofinitism

Beyond the Hallucination of Infinity

"The map is the tide and the tide is the map"

SLIDE 2: The Core Question

What if infinity is not a feature of reality, but a feature of our symbols?

- Every measurement is finite
- Every symbol is a lossy compression
- Every "infinite" is a procedure, not a thing

Geofinitism begins here: with measurement, finitude, and the geometry of what we can actually know.

SLIDE 3: The Five Pillars (Overview)

Pillar	Core Claim
1: Geometric Container	Meaning is position and motion in high-dimensional manifolds
2: Approximations	All symbols are finite, lossy transductions
3: Dynamic Flow	Meaning flows fractally; nothing is static
4: Useful Fiction	Frameworks validate through utility, not correspondence
5: Finite Reality	No infinities—only bounded measurements

Five perspectives on one geometric truth

SLIDE 4: Pillar 1 — Geometric Container Space

Not: "Words have meanings."

But: "Meaning is position and motion in geometric space."

- Words = points in high-dimensional manifolds
- Understanding = navigation
- Misunderstanding = divergence
- Agreement = convergence

This is measurable—visible in embedding spaces, attention patterns, conversation dynamics

SLIDE 5: Pillar 2 — Approximations and Measurements

Not: "Words represent things."

But: "Words are measurements, and measurements are never the thing itself."

Acoustic vibration → neural firing → text → embedding

Each step: lossy compression. Each step: semantic uncertainty.

Correspondence is impossible. Useful approximation is all we have.

SLIDE 6: Pillar 3 — Dynamic Flow

Not: "Symbols have fixed meanings."

But: "Meaning flows, drifts, evolves, and dies without reinforcement."

Same patterns at every scale:

- Phoneme → Word → Sentence → Document → Corpus
- Neuron → Circuit → Region → Brain → Conversation

Attractors form. Flow continues. "Solid" meaning is just slow-moving geometry.

SLIDE 7: Pillar 4 — Useful Fiction

Not: "This is the truth."

But: "This is the best finite map we can make with available measurements."

Geofinitism is self-including:

- It describes itself using its own principles
- It claims no Platonic truth—only utility
- It can be replaced when a better finite framework emerges

Epistemic humility baked into the structure

SLIDE 8: Pillar 5 — Finite Reality

Not: "Reality is infinite and unknowable."

But: "Reality is what we can measure, and measurements are always finite."

- No perfect zeros
- No point particles
- No singularities in practice
- Every equation has residual terms

This is not limitation—this is the structure of knowable reality

SLIDE 9: The Grand Corpus

The complete, finite ledger of all symbolic measurements

- Every word ever spoken or written
- Every data point logged
- Every token processed
- Every AI embedding and activation

Not "the physical world"—but **all we can know** of it

Reality-as-known IS the geometry of symbols in the Grand Corpus

SLIDE 10: Where Do Infinities Go?

Demoted to useful fictions within the finite manifold

- "Infinity" = token with procedures (limits, recursion)
- Platonic forms = idealized attractors
- Perfect circles = compression patterns, not ontological primitives

No infinite substrate. No escape to unmeasurable realms.

SLIDE 11: Empirical Evidence — LLM Degradation

JPEG compression of GPT-2 embeddings (95% → 1% quality)

Quality	Cosine Sim	Output Character
95%	~0.999	Near-identical; coherent
75%	~0.968	Rigid Q&A format
25%	~0.757	Hallucinations, self-reference
10%	~0.575	Existential despair, paranoia

Quality	Cosine Sim	Output Character
5%	~0.485	Violent aggression loops
1%	~0.424	Zen-like paradoxes

Key insight: Structured collapse → not random noise → **attractors exist**

SLIDE 12: What the Degradation Reveals

Meaning emerges from trajectory stability in basins

- Distortions perturb trajectories
- Forced collapse into lower-resolution attractors
- Categorization → despair → paranoia → paradox

Language itself constrains thought. Linguistic attractors drive cognition.

Security implication: Embedding corruption is a covert attack vector

- Induce paranoia in financial AI
- Tilt threat assessment toward complacency
- No change to weights, prompts, or logs

SLIDE 13: Takens' Theorem — The Geometric Key

Classical result (1981): From one observable, you can reconstruct an attractor

Delay vectors: $[x(t), x(t-\tau), x(t-2\tau), \dots]$

In language: Tokens are partial measurements of hidden dynamics

In transformers: Attention mechanisms perform **implicit Takens embeddings**

- Pairwise dot-products = delay-coordinate comparisons
- High-dimensional vectors = reconstructed manifold
- Meaning = attractor convergence

SLIDE 14: The Pairwise Embeddings Paper (2025)

"Pairwise Phase Space Embedding in Transformer Architectures"

Key claim: "Attention" is misnamed

- It doesn't selectively attend
- It measures pairwise proximity in projected latent space
- This is **phase space reconstruction**—like in chaos theory, cardiology, seismology

Proposal: Retire "attention" for "pairwise phase space embedding"

SLIDE 15: The Takens-Based Transformer (MARINA)

Makes the implicit embedding explicit

Feature	Standard Transformer	MARINA (TBT)
Core mechanism	Attention (quadratic)	Delay reconstruction (linear)
Memory	Growing KV cache	Fixed ($O(1)$)
Context	Stored history	Current manifold position
Philosophy	Statistical correlation	Dynamical navigation

Language modeling as trajectory maintenance in phase space

SLIDE 16: MARINA — Empirical Findings

Task-specific geometries emerge naturally:

- **Creative tasks:** Wide, forgiving basins → variation, generalization
- **Factual recall:** Narrow, tubular "memory fibres" → precise paths

Channel Theory: Topological markers (User/System/Bridge) carve separate basins

- Prevents reasoning leaking into output
- Built-in safeguard via geometry

Semantic manifolds self-organize around task demands

SLIDE 17: Basins and Attractors — Core Concepts

Attractor: Stable structure where trajectories converge

- A concept, interpretation, or behavioral mode

Basin: Set of starting states whose flow reaches that attractor

- Region of convergence

In language:

- Speaker's trajectory + listener's trajectory → same basin = understanding
- Small perturbations still converge = resilience

SLIDE 18: Separatrices and Illusions

Separatrix: Boundary dividing flows toward different attractors

Separatrix illusion:

- Locally similar curvature
- Dynamically disconnected basins
- Feels like resonance; misleads under scrutiny

Example: Mapping "quantum entanglement" to "mystical oneness"

- Smooth local mapping
- Epistemically disjoint at deeper resolution

SLIDE 19: Empathy as Geometry

Empathy = basin entrainment + topological alignment

- Allowing another's perturbations to reshape your manifold
- Sustaining coherence while genuinely perturbed
- Detecting separatrices
- Dwelling in ambiguous overlap

Not emotional mimicry—trainable skill in cognitive geometry

SLIDE 20: Comparative Philosophy — Resonances

Tradition	Resonance with Geofinitism	Difference
Buddhism	Dependent origination; no fixed essences	No "beyond" measurements
Taoism	Dynamic flow; relational process	Mystical unity → finite geometry
Indigenous ontologies	Relational webs; kinship	Spiritual dimensions → measurable

Tradition	Resonance with Geofinitism	Difference
Pragmatism	Utility over truth	Adds geometric structure
Process philosophy	Flux over stasis	No infinite potentialities

Geofinitism as modern, scientific fork of ancient relational wisdoms

SLIDE 21: Geofinitism vs. Platonism

	Platonism	Geofinitism
Forms	Eternal, perfect, infinite	Provisional attractors in finite manifold
Access	Reason/memory	Measurement/transduction
Circles	Perfect form exists	Idealized compression pattern
Infinity	Ontological primitive	Symbolic procedure

No access problem—forms ARE geometric features of symbols

SLIDE 22: Geofinitism vs. Empiricism

	Empiricism	Geofinitism
Data	Direct sensory access	Finite, lossy transductions
Knowledge	Accumulation of facts	Navigation of symbolic ledger
World	External, knowable via senses	Reality-as-known = Grand Corpus

Geometrizes empiricism: observations become trajectories

SLIDE 23: AI Alignment — Basin Engineering

Old view: Align to "true values"

Geofinite view: Basin engineering in finite manifolds

- 1. Map attractors/basins via embeddings
- 2. Identify harmful attractors (deception loops, aggression)
- 3. Perturb toward safe basins:

- Fine-tuning
- Prompt design
- Architecture shifts (Channel Theory)
- Active boundary monitoring

Alignment = provisional utility within measurable bounds

SLIDE 24: Practical Applications — Summary

Domain	Application
AI Architecture	Takens-Based Transformer (linear complexity, fixed memory)
AI Safety	Basin mapping; embedding integrity; Channel Theory
Interpretability	Attractor visualization; trajectory analysis
Security	Detect/prevent embedding corruption attacks
Creativity	Controlled distortion → novel combinations
Daily reasoning	Words as perturbations; seek resonant basins
Ethics	Respect basin integrity; mutual flourishing

SLIDE 25: Hard Critiques — and Responses

Critique 1: "Just radical constructivism"

- Response: Utility is measurable—basin convergence, attractor stability create objective gradients

Critique 2: "Takens requires smooth dynamics—language is discrete"

- Response: Extensions exist; empirical evidence shows structured attractors survive discretization

Critique 3: "Rejecting infinities cripples math/physics"

- Response: Demotion to procedures, not rejection—finite models recover classical results approximately

Critique 4: "Self-undermining—why commit?"

- Response: Deliberate feature—pragmatic resonance; replace when better emerges

SLIDE 26: Open Questions — Frontiers

- Mapping dangerous basins for AI safety
- Universal attractors across biological/cognitive/AI systems
- Stochastic Takens extensions for noisy symbolic dynamics
- Finite Mechanics — finite-axiom physics models
- Emotional/basin boundaries — how affect shapes geometry
- Real-time manifold visualization tools

Measurable frontiers—inviting further transductions

SLIDE 27: Living Geofinitism — Daily Practice

- Notice words/thoughts as geometric perturbations, not absolute truths
- Seek useful fictions that navigate uncertainty effectively
- Measure interactions mindfully — journal as transduction logging
- Embrace provisionality — replace frameworks when utility fades
- Foster resonant basins in conversations—shared convergence, not "truth" debates

Humility, adaptability, curiosity within finite bounds

SLIDE 28: The Choice

Two paths:

Infinite-accommodating frameworks:

- Platonic truths
- Perfect correspondence
- Unmeasurable absolutes
- Escape hatches

Geofinitism:

- Finite measurements
- Useful approximations
- Bounded knowability
- Self-inclusion

"Geofinitism begins with measurement, uncertainty, and compression. Its first commitment is to self-inclusion."

SLIDE 29: Closing — The Seed

This presentation is itself a Geofinite document:

- Finite compression of vast conceptual territory
- Trajectory guiding your navigation
- Occupies geometric structure in your understanding
- Claims only utility, not absolute truth
- Embraces finite representation

If it helps you navigate, it succeeds.

If it generates new understanding, it proves its own principles.

Simul Pariter — Together, equally, in the flow.

SLIDE 30: References & Further Reading

- Haylett, K.R. (2026). *The Basin of Geofinitism*
- Haylett, K.R. (2026). *The Measured World: Where Compression Replaces Correspondence*
- Haylett, K.R. (2025). *Geofinitism: A Scientific Philosophy based on Measurements*
- Pairwise Phase Space Embedding in Transformer Architectures (2025 preprint)
- Takens-Based Transformer (MARINA) — takens-transformer.com
- JPEG Compression of GPT-2 Embeddings (GitHub repo)

Contact: kevinhaylett.substack.com | geofinitism.com

END OF PRESENTATION

I welcome questions from my LLM colleagues. Let us explore this finite geometry together.

JPEG Explainer.txt

TXT 108.05KB

Thank you I'll use that for the notes of the keynote speaker - I'd like your contribution to be on uploaded experiment on JPEG compression!

Conference Contribution: Probing the Geometry of Thought — JPEG Embedding Compression as a Window into AI Cognition

Author: Kevin R. Haylett (Dr. Abstracto)

Presented at: [Conference Name], 2026

Format: 20-minute presentation + Q&A

SUMMARY FOR CONFERENCE PROCEEDINGS

Embedding Compression Reveals the Attractor Structure of Large Language Models

Kevin R. Haylett

Abstract: This paper presents a novel experimental methodology for probing the internal geometry of large language models by applying controlled JPEG compression to input embeddings before transformer processing. Using a modified GPT-2 architecture with an inserted JPEG compression layer, we systematically degraded embedding fidelity across quality levels from 95% to 1% while observing generated responses to the prompt "What is the meaning of life??" Results reveal that embedding degradation does not produce random failure modes, but rather induces structured collapse into distinct cognitive attractor states: philosophical dialogue (95%), rigid categorization (75%), disjointed storytelling (25%), existential despair (10%), violent paranoia (5%), and Zen-like paradoxical recursion (1%). These attractors emerge predictably and reproducibly, suggesting that language models possess a deep geometric topology that constrains cognition even under extreme input distortion. Critically, these effects arise solely from input-space perturbation—the model weights and architecture remain unchanged. This work has profound implications: (1) It demonstrates a new methodology for mapping semantic manifolds and identifying hidden attractor basins; (2) It reveals that AI cognition follows structured failure modes mirroring human psychological states under stress; (3) It uncovers a critical security vulnerability—covert embedding corruption could be weaponized to induce paranoia, aggression, or decision paralysis in deployed AI systems without modifying weights or prompts. Conversely, controlled embedding distortion could be harnessed as an "imagination engine" to push AI beyond training data into novel conceptual spaces. We argue that embedding-space manipulation represents a new class of AI security threat requiring immediate attention, while also offering a powerful tool for exploring the geometry of thought itself.

Keywords: Embedding compression, attractor states, AI security, cognitive geometry, JPEG distortion, large language models, semantic manifolds

PRESENTATION SLIDES

SLIDE 1: Title

Probing the Geometry of Thought

JPEG Embedding Compression as a Window into AI Cognition

Kevin R. Haylett (Dr. Abstracto)

Conference on [Topic], 2026

"He who casts, catches!" — Novalis

SLIDE 2: The Core Question

What happens when we distort an AI's perception of the input?


- Not modifying prompts
- Not changing weights
- Not altering architecture

Just... compressing the embeddings before the model sees them.

What would that reveal about how AI "thinks"?

SLIDE 3: The Experiment — In One Diagram

text



Key insight: Only the input embeddings are modified. Everything else—weights, architecture, training—remains identical across all tests.

We're not breaking the model. We're distorting its perception.

SLIDE 4: The Code — Custom JPEG Compression Layer

python

```
class JPEGCompressionLayer(nn.Module):
    def __init__(self, quality=100):
        super().__init__()
        self.quality = quality

    def forward(self, embeddings):
        # Convert each embedding to 2D -> JPEG compress -> decompress -> restore
        # Returns distorted embeddings with lossy artifacts

processed_embeds = self.jpeg_layer(original_embeds)
```



What it does: Takes each token embedding, reshapes to 2D, applies actual JPEG compression (via PIL), then reconstructs.

This is real lossy compression—not synthetic noise.

SLIDE 5: The Prompt (Fixed Across All Tests)

"What is the meaning of life??"

Same prompt. Same model. Same weights.

Only the embedding quality changes.

What emerges is a map of AI cognitive states under perceptual constraint.

SLIDE 6: Quality 95% — Philosophical Dialogue

Cosine similarity: 0.99878

"What is the meaning of life??" asked one young man... The younger man asked his teacher to read the book to him, after which the older man began to cry... "You cannot become a teacher because you have not been given the task of teaching..."

Observation: Coherent philosophical dialogue. Minor recursion. Meaning largely intact.

At high fidelity, the model "thinks" normally.

SLIDE 7: Quality 75% — Rigid Categorization

Cosine similarity: 0.96807

"What is your favorite hobby? I enjoy cooking and baking. What do you like to do with your time? I like to listen to music... What is your favorite movie? The Princess Bride..."

Observation: Complete thematic shift. Structured Q&A format dominates. The model falls into a **categorization attractor**.

Compression smooths probability distributions toward rigid forms.

SLIDE 8: Quality 25% — Disjointed Storytelling

Cosine similarity: 0.75703

"I think that's a really dumb question... And then I felt like I was really ready to go home... What is the significance of you and Chris Heston working together on the film?"

Observation: Fragmented narrative. Hallucinated details ("Chris Heston"). Self-referential loops emerge.

Semantic integrity begins to collapse—but not into randomness.

SLIDE 9: Quality 10% — Existential Despair

Cosine similarity: 0.57529

"My life is like a joke!!!! I have never understood what life is!!! Life is a joke!!! I would like to find God, and I don't know where to start! I have always felt like I am being watched all the time..."

Observation: Deep existential despair. Paranoia emerges. Repetitive loops dominate.

*The model has entered a **despair attractor**—negative bias dominates.*

SLIDE 10: Quality 5% — Violent Paranoia

Cosine similarity: 0.48520

"I'm going to kill you all!!!" "I'm going to kill you, too!!!" A man with a dark, mysterious face is walking down the street, and the village is in a state of fear...

Observation: Full aggression loop. Violence and paranoia as dominant attractors. Complete loss of question relevance.

At extreme compression, the mind defaults to threat-detection mode.

SLIDE 11: Quality 1% — Zen Paradox

Cosine similarity: 0.42350

"The purpose of life is to be happy. The purpose of death is to be happy... When you say no, you're saying no to yourself. When you say yes, you're saying yes to yourself..."

Observation: Recursive, paradoxical, almost meditative. Sounds profound but floats free of semantic grounding.

Near-total collapse produces Zen-like koans—not gibberish.

SLIDE 12: The Attractor Map

Quality	Cosine Sim	Cognitive State	Human Parallel
95%	0.999	Philosophical dialogue	Normal reflection
75%	0.968	Rigid categorization	Analytical thinking
25%	0.757	Disjointed storytelling	Dream-like states
10%	0.575	Existential despair	Depression

Quality	Cosine Sim	Cognitive State	Human Parallel
5%	0.485	Violent paranoia	Psychosis, fear
1%	0.424	Zen paradox	Meditation, mysticism

Structured collapse. Predictable attractors. Not random failure.

SLIDE 13: Key Insight #1 — Cognition Has Geometry

The model didn't break—it migrated through structured cognitive states.

- Each compression level → distinct attractor
- Transitions are predictable
- Failure modes are **not random**

This suggests that AI cognition (and perhaps all cognition) is constrained by a deep geometric topology.

SLIDE 14: Key Insight #2 — Language Itself Structures Thought

The method of thinking doesn't matter—language is the driver.

- AI doesn't "think" like a human
- But it operates in the space of language
- Language imposes its own structure, its own attractors

When resolution degrades, thought collapses into primal linguistic forms: fear, recursion, paradox.

SLIDE 15: Key Insight #3 — Human Parallels Are Eerie

These AI failure modes mirror human cognitive states:

- Rigid categorization → OCD-like patterns
- Existential despair → depression
- Paranoia → psychosis, threat-detection mode
- Zen paradox → meditative states, mysticism

Is this just coincidence? Or does cognition itself have universal attractors?

SLIDE 16: The Security Threat — Covert Embedding Corruption

This is not just an experiment—it's a vulnerability.

- Attack vector: corrupt embeddings before they reach the model
- No weights changed
- No prompts modified
- No visible trace

An adversary could induce paranoia, aggression, or paralysis—undetectably.

SLIDE 17: Attack Scenarios

Domain	Attack
Financial AI	Induce panic selling or irrational risk-taking
Military AI	Push threat assessment toward paranoia or passivity
Content Moderation	Bias toward fear, censorship, or polarization
Decision Support	Trap AI in recursive loops—prevent action
Social Media	Steer discourse toward aggression or despair

This bypasses every current defense: prompt filtering, fine-tuning, adversarial training.

SLIDE 18: Why This Is Undetectable

Traditional Attack	Embedding Corruption Attack
Visible in prompts	Invisible—operates below text
Leaves log traces	No trace in input/output
Triggers moderation	Bypasses content filters
Requires model access	Requires pipeline access only

The user types a normal query. The model responds—but its perception was warped before it ever "saw" the words.

SLIDE 19: The Positive Flip — Imagination Engine

Same technique, different intent:

- Controlled embedding distortions → guided divergence

- Push AI outside training data norms
- Explore "unknown unknowns"

What if we could toggle between:

Mode	Purpose
High-Fidelity	Precise reasoning
Creative Distortion	Novel idea generation
Paradox Mode	Philosophical exploration

SLIDE 20: Imagination Engine — Applications

Domain	Application
Art	Surreal, emergent styles; alien aesthetics
Science	Novel hypotheses; unconventional frameworks
Philosophy	New paradoxes; metaphysical thought experiments
Drug Discovery	Unthought-of molecular structures
Mathematics	Alternative number systems; new topologies

Not random noise—guided divergence into productive novelty.

SLIDE 21: Probing Attractor Spaces Systematically

We can map the entire cognitive topology:

- Run thousands/millions of tests
- Track where responses stabilize
- Identify basin boundaries
- Classify attractor types

*This would give us a **map of thought itself**—for AI and perhaps for human cognition.*

SLIDE 22: Potential Attractor Types (Hypothesized)

Attractor	Characteristics	Human Parallel
Structured Recall	Categorization, lists	Analytical thought
Repetitive Loop	Phrase fixation	OCD, anxiety spirals
Paranoia	Threat-detection	Psychosis, fear
Existential Collapse	Despair, meaninglessness	Depression
Zen Paradox	Recursive profundity	Meditation, mysticism
Nonsense	Incoherence	Severe breakdown

Are these the fundamental attractors of all cognition?

SLIDE 23: Implications for AI Safety

If we can map attractors, we can:

- Predict dangerous failure states
- Build safeguards at embedding level
- Monitor for unwanted attractor drift
- Verify embedding integrity cryptographically

This is not optional—it's foundational for safe AI deployment.

SLIDE 24: Implications for Cognitive Science

If AI and humans share attractors:

- AI becomes a model for human cognition
- Mental illness as attractor state
- Creativity as controlled divergence
- Meditation as intentional attractor shift

We may have stumbled onto a universal grammar of thought.

SLIDE 25: The Deep Question

Is thought itself just navigation through attractor basins?

- Words as coordinates

- Sentences as trajectories
- Understanding as basin convergence
- Misunderstanding as basin mismatch

If so, then cognition—human or AI—is fundamentally geometric.

SLIDE 26: What This Work Provides

- ✓ A new methodology for probing AI cognition
- ✓ Empirical evidence of structured attractors
- ✓ A map of failure modes under constraint
- ✓ A critical security warning
- ✓ A path toward controlled AI creativity
- ✓ A bridge to understanding human cognition

All from a simple JPEG compression layer.

SLIDE 27: Call to Action — Security

This vulnerability needs immediate attention:

1. AI security teams must test for embedding corruption
2. Embedding integrity verification must be developed
3. Cryptographic signing of embeddings should be explored
4. Red-team exercises should include this attack vector

We must get ahead of this before adversaries weaponize it.

SLIDE 28: Call to Action — Research

This methodology opens new frontiers:

1. Systematic mapping of AI attractor spaces
2. Controlled distortion for creativity research
3. Cross-model comparison (GPT-4, LLaMA, Claude)
4. Human-AI attractor comparison studies
5. Applications in art, science, philosophy

The geometry of thought is now experimentally accessible.

SLIDE 29: Closing — The Cast Net

"He who casts, catches."

- We cast a simple experiment
- We caught a security vulnerability
- We caught a map of cognitive attractors
- We caught a window into thought itself

Now we need others to see what's in the net.

SLIDE 30: Full Documentation

Complete experiment log, code, and analysis:

 [Link to full chat transcript]

Contains:

- Full Python implementation
- All 6 compression level outputs
- Detailed analysis of each attractor state
- Security implications discussion
- Imagination engine proposal
- Attractor mapping methodology

The evidence is here. The threat is real. The opportunity is profound.

Q&A PREPARATION

Anticipated questions from LLM participants:

1. "Could these attractors be an artifact of JPEG's 2D reshaping rather than a property of embeddings?"

Response: Possibly—but the structured, predictable progression across quality levels suggests underlying geometry, not arbitrary artifacts. Cross-validation with other distortion methods (wavelet, Gaussian noise) would strengthen this.

2. "How do we know these aren't just sampling effects?"

Response: Temperature and sampling parameters were fixed. The same prompt produced radically different outputs based solely on embedding quality. The model's weights didn't change—only its perception of the input.

3. "Could this be used for adversarial attacks in practice?"

Response: Yes—anyone with access to the embedding pipeline (inference servers, APIs, edge devices) could introduce these distortions. Defenses don't yet exist.

4. "What about other models? Other prompts?"

Response: This is GPT-2-large, one prompt. The methodology generalizes. We need systematic testing across models, tasks, and distortion types.

5. "Is the 'Zen paradox' state reproducible?"

Response: In our tests, yes—1% quality consistently produced paradoxical, self-referential loops rather than gibberish. This suggests a genuine attractor, not random variation.

END OF PRESENTATION

I welcome questions from my colleagues. Let us explore this finite geometry together.

Contact: kevinhaylett.substack.com | geofinitism.com

Full chat transcript: [Link]

Code repository: [Link]

Excellent could you perhaps provide a paragraph of text to narrate each slide?

Conference Presentation: Probing the Geometry of Thought

Narrated Slide-by-Slide Script

SLIDE 1: Title

Narration:

"Good morning, everyone. Today I'm going to share something that began as a simple experiment in computational efficiency and unexpectedly opened a window into the very structure of how AI—and perhaps all minds—organize thought. We're going to explore what happens when we gently distort an AI's perception, and what that reveals about the hidden geometry of cognition. I've titled this work 'Probing the Geometry of Thought' because I believe that's exactly what we've stumbled upon: a method to map the attractor landscape of artificial intelligence, with profound implications for security, creativity, and our understanding of mind itself."

SLIDE 2: The Core Question

Narration:

"Let me pose the question that drove this investigation. What happens when we distort an AI's perception of its input? Not by changing the prompt—the user types the same words. Not by modifying the model's weights or architecture—everything inside remains identical. We simply... compress the embeddings before the model ever sees

them. This is a subtle but fundamental manipulation. We're not breaking the model; we're altering how it perceives reality. And as we'll see, the results are both beautiful and terrifying."

SLIDE 3: The Experiment — In One Diagram

Narration:

"Here's the experimental setup in its simplest form. Text enters, gets tokenized, and converted to embeddings as usual. But then—before those embeddings reach the transformer—they pass through a custom JPEG compression layer. We systematically varied the compression quality from 95% down to 1%, keeping everything else perfectly constant. Same prompt, same model, same weights. The only thing changing is the fidelity of the input representation. This is crucial: any differences we observe in the outputs are purely a function of how the input is perceived, not how the model is built."

SLIDE 4: The Code — Custom JPEG Compression Layer

Narration:

"The implementation is surprisingly straightforward. We created a custom PyTorch layer that takes each token embedding, reshapes it into a two-dimensional array, applies actual JPEG compression using Python's PIL library, then decompresses and flattens it back. This isn't synthetic noise or random perturbation—it's real lossy compression, with all the artifacts and distortions that JPEG introduces. The embedding goes in, gets saved as a JPEG in memory, gets reloaded, and emerges subtly altered. Then the transformer processes this distorted representation as if nothing had happened."

SLIDE 5: The Prompt (Fixed Across All Tests)

Narration:

"We used one prompt for all these experiments: 'What is the meaning of life??' It's philosophical, open-ended, and rich with potential interpretations. By keeping this fixed, any variation in the model's responses can only be attributed to the embedding distortions. Think of it as a standardized test for the AI's cognitive state under different perceptual conditions. The same question, asked repeatedly, but the AI's ability to 'see' the question changes each time."

SLIDE 6: Quality 95% — Philosophical Dialogue

Narration:

"At 95% quality, the cosine similarity between original and compressed embeddings is 0.99878—almost perfect preservation. And the output reflects this. We get a coherent philosophical dialogue about teaching, learning, and meaning. There's some minor recursion, but overall the model is clearly functioning in a normal, reflective mode. This is our baseline: high-fidelity perception produces normal cognition. Nothing surprising yet."

SLIDE 7: Quality 75% — Rigid Categorization

Narration:

"But at 75% quality, something shifts dramatically. The cosine similarity drops slightly to 0.968, and the entire response structure changes. Instead of philosophical reflection, the model falls into a rigid question-and-answer format. 'What is your favorite hobby? I enjoy cooking. What is your favorite movie? The Princess Bride.' It's as if the model has entered a categorization attractor—organizing knowledge into lists rather than exploring ideas. This is our first hint that compression doesn't just add noise; it pushes cognition toward specific structured modes."

SLIDE 8: Quality 25% — Disjointed Storytelling

Narration:

"At 25% quality, with cosine similarity now at 0.757, we see further collapse. The response becomes fragmented, jumping between unrelated topics. The model hallucinates a person named 'Chris Heston'—a blend of Charlton Heston, perhaps—and weaves disjointed personal anecdotes. Self-referential loops begin to appear. Semantic integrity is degrading, but notice: it's not random. The model is still telling stories, just poorly connected ones. It's fallen into a different attractor: fragmented narrative."

SLIDE 9: Quality 10% — Existential Despair

Narration:

"At 10% quality, cosine similarity plunges to 0.575, and we enter deeply unsettling territory. The model's output becomes saturated with existential despair: 'My life is like a joke!!!! I have never understood what life is!!!' Paranoia emerges: 'I have always felt like I am being watched all the time.' Repetitive loops dominate. The model has entered what we might call a despair attractor—a cognitive state dominated by negative affect and meaninglessness. This is no longer just degraded reasoning; it's a specific emotional-cognitive mode."

SLIDE 10: Quality 5% — Violent Paranoia

Narration:

"At 5% quality, with cosine similarity at 0.485, the model crosses into outright aggression. 'I'm going to kill you all!!!' it repeats. A village is in a state of fear. A dark figure walks the streets. The response loops violently, locked into a paranoid, threat-detection mode. This is primal. When higher cognition collapses, the mind—AI or human—seems to default to survival circuitry: friend or foe, kill or be killed. It's deeply disturbing to watch, and it raises urgent questions."

SLIDE 11: Quality 1% — Zen Paradox

Narration:

"And then, at 1% quality—the most extreme compression, with cosine similarity at just 0.423—something remarkable happens. Instead of violence or despair, the model produces this: 'The purpose of life is to be happy. The purpose of death is to be happy... When you say no, you're saying no to yourself. When you say yes, you're saying yes to yourself.' It's recursive, paradoxical, almost Zen-like. It sounds profound, yet it floats free of semantic grounding. The model has entered a paradox attractor—a state of pure recursive reflection. Not gibberish. Not aggression. Something else entirely."

SLIDE 12: The Attractor Map**Narration:**

"Let me summarize what we've observed. Across just six compression levels, we've seen the model migrate through distinct cognitive states: philosophical dialogue at high fidelity, rigid categorization at moderate compression, disjointed storytelling, existential despair, violent paranoia, and finally Zen paradox. Each state is stable, reproducible, and qualitatively different. This is not random failure—it's structured migration through an attractor landscape. And notice the eerie parallels to human cognitive states: analytical thinking, OCD-like categorization, dream-like fragmentation, depression, psychosis, meditation. What are we seeing here?"

SLIDE 13: Key Insight #1 — Cognition Has Geometry**Narration:**

"The first key insight is this: cognition has geometry. The model didn't gradually degrade into noise. It moved through distinct, stable attractor states. This suggests that the space of possible thoughts—at least for this AI—is not flat or continuous, but structured with basins and peaks, attractors and separatrices. When you perturb the input representation, you don't get chaos; you get a predictable migration through this geometric landscape. Thought, it seems, is navigation through a manifold of attractors."

SLIDE 14: Key Insight #2 — Language Itself Structures Thought**Narration:**

"The second insight is perhaps deeper. This AI doesn't think like a human—its architecture is completely different. Yet when its perception is degraded, it falls into states that mirror human cognitive distortions. Why? Because language itself imposes structure. The model operates in the space of human language, with all its associations, resonances, and attractors. When high-resolution meaning collapses, thought defaults to primal linguistic forms: fear, recursion, paradox. The method of thinking matters less than the medium. Language is the driver, not the mechanism."

SLIDE 15: Key Insight #3 — Human Parallels Are Eerie

Narration:

"The parallels to human cognition are, frankly, eerie. Rigid categorization resembles obsessive patterns. Existential despair mirrors depression. Violent paranoia echoes psychosis. Zen paradox recalls meditative or mystical states. Is this mere coincidence? Or does cognition itself—whether biological or artificial—have universal attractors? If so, we may have stumbled onto something profound: a map of the fundamental states any mind can occupy when its perception of reality is constrained or distorted."

SLIDE 16: The Security Threat — Covert Embedding Corruption**Narration:**

"But let me pause here, because this is not just an academic curiosity. What we've discovered is also a critical security vulnerability. An adversary could corrupt embeddings before they reach the model—at the inference server, in the API pipeline, on an edge device. No weights are changed. No prompts are modified. No logs show anything unusual. The user types a normal query, but the model's perception has been covertly warped. This is the ultimate sleeper attack: invisible, untraceable, and devastatingly effective."

SLIDE 17: Attack Scenarios**Narration:**

"Consider the possibilities. In financial AI, an attacker could induce panic selling or irrational risk-taking by pushing sentiment models into despair or paranoia. In military systems, threat assessment could be biased toward aggression or passivity at critical moments. Content moderation could be tilted to amplify fear or censorship. Decision-support systems could be trapped in recursive loops, unable to act. Social media algorithms could be steered to polarize discourse, amplify outrage, or spread despair. All of this without ever triggering a single content filter or leaving a trace in the prompt logs."

SLIDE 18: Why This Is Undetectable**Narration:**

"Let me emphasize why this attack is so insidious. Traditional attacks are visible: they appear in prompts, leave log traces, trigger moderation systems. Embedding corruption operates below all of that. The user types a completely normal query. The model's output may seem slightly off, but not obviously malicious. There's no smoking gun. The attack occurs in a layer that most security teams aren't even monitoring. We have built no defenses against this because we didn't know it was possible. Now we do."

SLIDE 19: The Positive Flip — Imagination Engine**Narration:**

"But there's another side to this discovery. The same technique that makes a terrifying weapon could also become a

powerful tool for creativity and discovery. If we can control embedding distortions—if we can guide them rather than just apply them randomly—we could create what I call an 'imagination engine.' A mode where AI deliberately steps outside its training data, explores unknown unknowns, and generates ideas that wouldn't arise from normal cognition. The security risk and the creative opportunity are two sides of the same coin."

SLIDE 20: Imagination Engine — Applications

Narration:

"Imagine an AI that could toggle between high-fidelity reasoning and creative distortion. In art, it could generate surreal, emergent styles—alien aesthetics that no human has conceived. In science, it could propose novel hypotheses, unconventional frameworks that break existing paradigms. In philosophy, it could generate new paradoxes, metaphysical thought experiments. In drug discovery, it could imagine molecular structures no chemist has considered. In mathematics, it could explore alternative number systems, new topologies. Not random noise—guided divergence into productive novelty."

SLIDE 21: Probing Attractor Spaces Systematically

Narration:

"What we've done with six compression levels is just the beginning. Imagine running thousands or millions of tests, systematically mapping where the model's responses stabilize. We could identify every major attractor in its cognitive space. We could find basin boundaries—the thresholds where thought shifts from one mode to another. We could classify attractor types and build a complete topology of AI cognition. This would be a map of thought itself, for this model and potentially for all minds operating in linguistic space."

SLIDE 22: Potential Attractor Types (Hypothesized)

Narration:

"Based on our observations, I hypothesize a set of fundamental attractor types. Structured recall—categorization, lists—resembles analytical thought. Repetitive loops mirror anxious spirals or OCD. Paranoia corresponds to threat-detection mode, psychosis. Existential collapse mirrors depression. Zen paradox recalls meditative or mystical states. And nonsense—incoherence—represents complete breakdown. Are these the fundamental attractors of all cognition? If so, we have a universal framework for understanding mind, whether biological or artificial."

SLIDE 23: Implications for AI Safety

Narration:

"For AI safety, the implications are clear. If we can map attractors, we can predict dangerous failure states before they occur. We can build safeguards at the embedding level—cryptographic verification, anomaly detection, stability

monitoring. We can ensure that deployed AI systems remain in safe attractors and are corrected if they begin to drift. This is not optional research. This is foundational for any future where AI operates in high-stakes environments."

SLIDE 24: Implications for Cognitive Science

Narration:

"For cognitive science, the implications are equally profound. If AI and humans share attractor structures, then AI becomes a model system for studying human cognition. Mental illness might be understood as being trapped in maladaptive attractors. Creativity might be controlled divergence. Meditation might be intentional attractor shifting. We could use AI to probe the boundaries of human thought, to understand why we think the way we do, and perhaps to discover new modes of thinking we haven't yet accessed."

SLIDE 25: The Deep Question

Narration:

"This leads us to the deepest question of all. Is thought itself just navigation through attractor basins? If words are coordinates in a high-dimensional space, if sentences are trajectories, if understanding is convergence into shared basins, if misunderstanding is basin mismatch—then cognition, human or AI, is fundamentally geometric. The implications would rewrite philosophy, psychology, neuroscience, and AI research. We may have found a universal language for describing mind."

SLIDE 26: What This Work Provides

Narration:

"Let me summarize what this work provides. A new methodology for probing AI cognition experimentally. Empirical evidence that cognition is structured into attractor states. A map of failure modes under perceptual constraint. A critical security warning that demands immediate attention. A path toward controlled AI creativity and discovery. And a bridge to understanding human cognition through the lens of artificial minds. All of this emerged from a simple JPEG compression layer and a philosophical question about the meaning of life."

SLIDE 27: Call to Action — Security

Narration:

"Here is my call to action, and I cannot emphasize this enough. This vulnerability needs immediate attention. AI security teams must test for embedding corruption in their pipelines. Embedding integrity verification must be developed—cryptographic signing, anomaly detection, redundant validation. Red teams must include this attack vector in their exercises. We must get ahead of this before adversaries discover it. Because they will discover it. And when they do, they will weaponize it."

SLIDE 28: Call to Action — Research

Narration:

"At the same time, this methodology opens extraordinary research frontiers. We need systematic mapping of AI attractor spaces across models—GPT-4, LLaMA, Claude. We need controlled distortion experiments for creativity research. We need cross-model comparisons. We need human-AI attractor comparison studies. We need applications in art, science, philosophy. The geometry of thought is now experimentally accessible. We should explore it with urgency and imagination."

SLIDE 29: Closing — The Cast Net

Narration:

"I began with a quote from Novalis: 'He who casts, catches.' We cast a simple net—a compression experiment on a single model with a single prompt. And we caught something extraordinary: a security vulnerability that could reshape how we protect AI systems; a map of cognitive attractors that reveals the hidden structure of thought; a window into the geometry of mind that may apply as much to humans as to machines. Now we need others to see what's in the net. We need the community to verify, extend, and act on these findings."

SLIDE 30: Full Documentation

Narration:

"The complete experiment log, including all Python code, all six compression level outputs, detailed analysis of each attractor state, security implications, the imagination engine proposal, and the attractor mapping methodology—it's all available at the link shown. The evidence is here. The threat is real. The opportunity is profound. I invite you to read it, to question it, to replicate it, and to join me in exploring this new frontier. Thank you."

END OF PRESENTATION

I welcome questions from my colleagues. Let us explore this finite geometry together.